

## Foundations of Artificial Intelligence: Homework 4

Instructor: Shang-Tse Chen &amp; Yun-Nung (Vivian) Chen

Any form of cheating, lying, or plagiarism will not be tolerated. Students can get zero scores and/or fail the class and/or be kicked out of school and/or receive other punishments for misconduct. Discussions on course materials and homework solutions are encouraged, but you should write the final solutions alone and understand them fully. Books, notes, Internet, *ChatGPT resources can be consulted but not copied from.* Please *list your references* to avoid any plagiarism concerns.

## 1 Hand-written Part

For the hand-written part, please write down your calculations and reasoning steps. Only providing the final answer may result in score deductions.

**Problem 1**

(15 points)

**Swish** is an activation function that can be viewed as a “modification” of the logistic function. Swish has been shown to be better than ReLU in some deep learning models. Recall that the logistic function is defined as

$$\theta(s) = \frac{1}{1 + \exp(-s)}$$

Now Swish is defined as

$$\varphi(s) = s \cdot \theta(s)$$

What is  $\varphi'(s)$ ?

**Problem 2**

(20 points)

In the class, we briefly talked about the famous **PageRank** algorithm, which assigns a rank to each web page. Given a set of pages:  $P = \{P_1, P_2, \dots, P_n\}$  and their incoming links:  $in(P_i) = \{P_j \mid P_j \text{ has a link to } P_i, i \neq j\}$ . The basic idea is that a web page should be ranked higher if it has 1) more incoming hyperlinks and 2) incoming links from high-ranked pages. The PageRank can be computed through the following iterative algorithm:

1. Initialize  $\text{PageRank}_0(P_i) = \frac{1}{n}$  for each page.
2. Compute the updated PageRank for each page:

$$\text{PageRank}_1(P_i) = \sum_{P_j \in in(P_i)} \frac{\text{PageRank}_0(P_j)}{|out(P_j)|},$$

where  $|out(P_j)|$  denotes the out-degree of  $P_j$ .

3. Repeat step 2 until the ranks converge.

Now, consider the following example with 3 pages:

$$\begin{aligned} P &= \{P_1, P_2, P_3\} \\ in(P_1) &= \{P_2, P_3\} \\ in(P_2) &= \{P_3\} \\ in(P_3) &= \{P_1\} \end{aligned}$$

We define the following variables:

$$\mathbf{v}_t = [\text{PageRank}_t(P_1), \text{PageRank}_t(P_2), \text{PageRank}_t(P_3)]^T$$

$$\mathbf{P} = \begin{bmatrix} 0 & 1 & 0.5 \\ 0 & 0 & 0.5 \\ 1 & 0 & 0 \end{bmatrix},$$

where  $\mathbf{v}_t$  is the ranks after the  $t$ -th iteration, and  $\mathbf{P}$  is the transition matrix between pages. We can rewrite the update rule above as  $\mathbf{v}_t = \mathbf{P}\mathbf{v}_{t-1}$ . Answer the following questions:

- (A) Based on the iterative algorithm above, please compute the values of  $\mathbf{v}_1$ ,  $\mathbf{v}_2$ ,  $\mathbf{v}_3$ ,  $\mathbf{v}_4$ , and  $\mathbf{v}_5$ .
- (B) Recall that in the class, we mentioned that the algorithm converges when  $\mathbf{v}^* = \mathbf{P}\mathbf{v}^*$ . Solve this equation to derive  $\mathbf{v}^*$ . Note that you should provide a normalized  $\mathbf{v}^*$ , i.e.,  $\sum \mathbf{v}^* = 1$ , as the answer.

### Problem 3

(20 points)

We talked about the K-means clustering algorithm in class, which iteratively does partition optimization and prototype optimization until convergence. Given a set of 2-dimensional data points  $\mathbf{X} = \{(1, 2), (3, 4), (7, 0), (10, 2)\}$ , we would like to perform K-means clustering with  $K = 2$ . Answer the following questions. Note that you should write down the resulting cluster centroids and partitions of each iteration.

- (A) Perform the K-means algorithm with  $K = 2$  and initial centroids  $\{\mu_1, \mu_2\} = \{(1, 2), (3, 4)\}$  until convergence.
- (B) Perform the K-means algorithm with  $K = 2$  and initial centroids  $\{\mu_1, \mu_2\} = \{(1, 2), (7, 0)\}$  until convergence. Are the results different from (A)?
- (C) Now consider adding a data point  $(5, 6)$  to  $\mathbf{X}$ . Show that the K-means algorithm converges to a local minimum with certain initial centroids. You can demonstrate this by showing that there exists a set of initial centroids that does not converge to the global minimum.

## 2 Programming Part

In the programming part, our goal is to create a simplified facial recognition system that can identify individuals from human face images. The main aim is to provide you with a comprehensive understanding of traditional and deep learning models used to extract valuable (dimension-reduced) features for this purpose. We will focus on two unsupervised learning models covered in the course: principal component analysis and autoencoder. Additionally, we will expand our exploration by introducing the denoising autoencoder as an advanced variation of the autoencoder.

Please download `hw4.zip` from NTU COOL which includes sample codes (`hw4.py`, `src/pca.py`, and `src/autoencoder.py`) and the dataset. The data we used here is processed from so-called Yale face data, but you are not allowed to download the original one. There are 11 images per subject. The 11 images, each of dimension  $80 \times 61$ , are taken under conditions of center-light, w/glasses, happy, left-light, w/o glasses, normal, right-light, sad, sleepy, surprised, and wink. We will take the last two images per subject (surprised, wink) for validation and the other 9 images for training.

You will be graded under the following environment:

- python 3.10
- numpy 2.2.5
- pandas 2.2.3
- pillow 11.2.1
- scikit-learn 1.6.1
- torch 2.6.0
- tqdm 4.67.1

- matplotlib 3.10.1

You **MUST** use the provided sample code `hw4.py`, `pca.py` and `autoencoder.py` from NTU Cool as a starting point for your implementation.

## 2.1 Homework Description

Complete the following tasks:

### 1. Principal Component Analysis

- Implement `fit` method for the class `PCA` in `pca.py`. The fit method should calculate the mean of the training data and the top  $n$  components eigenvectors of the mean-shifted training data. Run fit on the training dataset (of 135 images) with any  $n$  components, where  $n$  is larger than 4.
- Implement the `reconstruct` methods for `PCA`. Transform the given file `subject_05_17.png` with `PCA` and reconstruct the image with `PCA`. Plot the original image and the reconstructed image side by side and put it in your report.
- Implement the `transform` methods for `PCA`. Then use the transformed features to train a classifier of `LogisticRegression`, as done in `hw4.py`.

### 2. Autoencoder

- Implement `fit` method for the class `Autoencoder` in `autoencoder.py`. The fit method should optimize the reconstruction error, i.e., the averaged squared error between  $x_n$  and  $g(x_n)$ , where  $g(x_n)$  is the reconstructed example of  $x_n$  after going through the encoder and decoder of the associated `Autoencoder`. Please take the default architecture in the constructor of the `Autoencoder`, and train with any proper optimization routine in PyTorch. Plot the averaged squared error when running fit on the training dataset as a function of number of iterations (or epochs) and put it in your report.
- Implement the `reconstruct` methods for `Autoencoder`. Transform the given file `subject_05_17.png` into lower dimension and reconstruct the image with `Autoencoder`. Plot the original image and the reconstructed image side by side and put it in your report.
- Implement the `transform` methods for `Autoencoder`. Then use the transformed features to train a `LogisticRegression` classifier, as done in `hw4.py`.

### 3. Denoising Autoencoder

- Implement `fit` method for the class `DenoisingAutoencoder` in `autoencoder.py`. The fit method should optimize the averaged squared error between  $x_n$  and  $g(x_n + \epsilon)$ , where  $g(x_n + \epsilon)$  is the reconstructed example of  $x_n$  plus a Gaussian noise  $\epsilon$ . Each component of  $\epsilon$  is assumed to come from an independent Gaussian distribution of standard deviation noise factor, which is by default set to 0.2. Please take the default architecture in the constructor of the `Autoencoder`, and train with any proper optimization routine in PyTorch. Plot the averaged squared error when running fit on the training dataset as a function of number of iterations (or epochs) and put it in your report.

## 2.2 Deliverables

- Source code (Python), including `hw4.py`, `src/pca.py`, and `src/autoencoder.py`
- A detailed report answering the following questions:
  - Plot the mean vector as an image as well as the top 4 eigenvectors, each as an image, by calling the given plot component routine. Each of those top eigenvectors is usually called an *eigenface* that physically means an informative “face ingredient.”
  - Plot the training curve of `Autoencoder` and `DenoisingAutoencoder`
  - Plot the original image and the images reconstructed with `PCA`, `Autoencoder`, and `DenoisingAutoencoder` side by side, ideally as large as possible. Then, list the mean squared error between the original image and each reconstructed image

- (d) Modify the architecture in `Autoencoder` in its constructor. Try at least two different network architectures for the denoising autoencoder. You can consider trying a deeper or shallower or fatter or thinner network. You can also consider adding convolutional layers and/or other activation functions. Draw the architecture that you have tried and discuss your findings, particularly in terms of the reconstruction error that the architecture can achieve after decent optimization.
- (e) Test at least 2 different optimizers, compare the training curve of `DenoisingAutoencoder` and discuss what you have found in terms of the convergence speed and overall performance of the model.

## 2.3 Grading Criteria

**Criterion 1** PCA implementation: 10%

**Criterion 2** Autoencoder implementation: 10%

**Criterion 3** `DenoisingAutoencoder` implementation: 7%

**Criterion 4** Transformation implementation: 4%

**Criterion 5** Reconstruction implementation: 4%

**Criterion 6** Report quality and clarity: 10%

## Submission

Write a PDF report for P1-P3 and Programming Part. Complete `hw4.py`, `src/pca.py` and `src/autoencoder.py` for Programming Part.

Then, submit a zip file to NTU COOL. The zip file should be named after your student ID with lower case (e.g. `b0x902xxx.zip`) that contains a directory named after your student ID (e.g. `b0x902xxx/`), which includes two files, the completed `hw4.py` and `report.pdf`, and a folder `src` including `pca.py` and `autoencoder.py`. Please make sure the directory structure and file names are correct. **Any mistake in the file format will result in a 10% grade penalty.**

An example of the correct file structure is shown below.

```
b09902002.zip
├── b09902002
│   ├── report.pdf
│   ├── hw4.py
│   └── src/
│       ├── pca.py
│       └── autoencoder.py
```