

# CheatSheet for Foundation of AI (midterm)

by Tawmial@NTU

## Adversarial Search

### Minimax Search (with alpha-beta pruning)

$\alpha$ : MAX's best option on path to root  
 $\beta$ : MIN's best option on path to root

```
def max-value(state,  $\alpha$ ,  $\beta$ ):
    initialize v = - $\infty$ 
    for successor of state:
        v = max(v, min-value(successor,  $\alpha$ ,  $\beta$ ))
        if v >  $\beta$ : return v
         $\alpha$  = max( $\alpha$ , v)
    return v
```

```
def min-value(state,  $\alpha$ ,  $\beta$ ):
    initialize v =  $+\infty$ 
    for successor of state:
        v = min(v, max-value(successor,  $\alpha$ ,  $\beta$ ))
        if v <  $\alpha$ : return v
         $\beta$  = min( $\beta$ , v)
    return v
```

### Depth-limited Search

- search only to a limited path in the tree
- use evaluation function for non-terminal positions

### Expectimax Search

- compute the average score under optimal play
- replace Min node with Chance node, where the outcome is uncertain
- perform better when against a random adversary

### Monte Carlo Tree Search (MCTS)

- do N rollout (simple, fast policy) from each child of the root, and count win-rate
  - pick the child with the best win-rate, only explore on that subtree
- MCTS 2.0: UCB for Tree**
- repeat until out of time:
    - given current search tree, recursively apply UCB to choose a path down to a leaf node n
    - add a new child c to n, and run a rollout from c
    - update the win counts from c back up to the root
  - choose the action leading to the child with highest N (# rollouts)

## Propositional Logic

- agent learn from a knowledge base (set of sentences in a formal language)
- agent can then inference new facts from the known facts (KB)
- can be used for localization (where am i), mapping (what's the environment like), planning

### Logic

- Syntax: what sentences are allowed?
  - World/Model: an assignment of literals
  - Semantics: which sentences are true in which worlds?
- |              |                   |
|--------------|-------------------|
| P $\wedge$ Q | {P=true, Q=false} |
| P $\wedge$ Q | means false       |

### Propositional Logic Syntax

- $\neg\alpha$  is true iff  $\alpha$  is false
- $\alpha \wedge \beta$  is true iff both  $\alpha$  and  $\beta$  are true
- $\alpha \vee \beta$  is true iff at least one in  $\alpha$  or  $\beta$  is true
- $\alpha \Rightarrow \beta$  is true iff  $\alpha$  is false or  $\beta$  is true
- $\alpha \Leftrightarrow \beta$  is true iff  $\alpha$  and  $\beta$  have the same value

### Inference

- Entailment:  $\alpha \models \beta$  iff in every world where  $\alpha$  is true,  $\beta$  is also true
- Proofs: an process to demonstrate entailment
  - model-checking: loop through every possible world
  - theorem-proving: use inference rules to derive new sentences

### Forward chaining

```
function FORWARD-CHAINING(KB, q)
  count ← a table, where count[c] is the number of symbols in c's premise
  inferred ← a table, where inferred[s] is initially false for all s
  agenda ← a queue of symbols, initially symbols known to be true in KB
  while agenda is not empty do
    p ← Pop(agenda)
    if p = q then return true
    if inferred[p] = false then
      inferred[p] ← true
      for clause c in KB where p is in c.premise do
        decrement count[c]
        if count[c] = 0 then add c.conclusion to agenda
  return false
```

### CNF

- a conjunction of clauses, each clause is a disjunction of literals
- distributivity:
  - $(A \vee B) \wedge C = (A \wedge C) \vee (B \wedge C)$
  - $(A \wedge B) \vee C = (A \vee C) \wedge (B \vee C)$
  - $\neg(A \wedge B) = (\neg A \vee \neg B)$
  - $\neg(A \vee B) = (\neg A \wedge \neg B)$

### Satisfiability

- a sentence is **valid** if it is true in every world
- a sentence is **satisfiable** if it is true in at least one world

### DPLL

- pure symbol: always appears with same "sign" in all clauses
- unit clause: contains exactly one unassigned symbol

```
function DPLL(clauses, symbols, partmodel={})
  if every clause in clauses is true in partmodel then return true
  if some clause in clauses is false in partmodel then return false
  clauses ← every clause in clauses that is not true (or false)

  P, value ← FIND-PURE-SYMBOL(symbols, clauses, partmodel)
  if P is non-null then
    return DPLL(clauses, symbols - {P}, partmodel U {P=value})

  P, value ← FIND-UNIT-CLAUSE(clauses, partmodel)
  if P is non-null then
    return DPLL(clauses, symbols - {P}, partmodel U {P=value})

  P ← First(symbols); rest ← Rest(symbols)
  return (DPLL(clauses, rest, partmodel U {P=true}) or
         DPLL(clauses, rest, partmodel U {P=false}))
```

## Unified Search

### Search Tree

- each node is a plan
- fringe = nodes under consideration
- decide which fringe nodes to expand

- Complete: Guarantee to find a solution if exists
- Optimal: Guarantee to find the least cost path

### Depth-First Search

- expand deepest node first
- fringe is a LIFO stack
- $O(b^d m)$  time;  $O(bm)$  space
- complete if no cycle; not optimal

### Breadth-First Search

- expand shallowest node first
- Fringe is a FIFO queue
- $O(b^d)$  time;  $O(b^d)$  space
- complete; optimal if costs are 1

### Uniform-Cost Search

- expand cost less node first
- effective depth  $d = C^*/e$ , where  $C^*$  is solution cost, e is min arc cost
- $O(b^d d)$  time,  $O(b^d)$  space
- complete and optimal

## Informed Search

### Heuristics function

- Admissibility:  $h(x) \leq \text{cost}(x \text{ to Goal})$
- Consistency:  $h(x - y) \leq \text{cost}(x \text{ to } y)$

### Relaxed Problem

- a simplified version of original problem by remove some constraints or add new operations
- solution can be heuristic for original problem

### A\* Search

- $f(n) = \text{path cost} + \text{goal proximity} = g(n) + h(n)$
- expand node with min  $f(n)$
- optimal if heuristic admissible (tree), consistent (graph)

```
function TREE-SEARCH(problem, fringe)
  fringe ← INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
  loop do
    if fringe is empty then return failure
    node ← REMOVE-FRONT(fringe)
    if GOAL-TEST([problem], STATE[node]) then return node
    for child-node in EXPAND(STATE[node], problem) do
      fringe ← INSERT(child-node, fringe)
```

```
function GRAPH-SEARCH(problem, fringe)
  closed ← an empty set
  fringe ← INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
  loop do
    if fringe is empty then return failure
    node ← REMOVE-FRONT(fringe)
    if GOAL-TEST([problem], STATE[node]) then return node
    if STATE[node] is not in closed then
      add STATE[node] to closed
      for child-node in EXPAND(STATE[node], problem) do
        fringe ← INSERT(child-node, fringe)
```

## Constraint Satisfaction Problem (CSP)

- state is defined by variables (each city)
- each variable has a domain (red, green, blue)
- constraints specify allowable assignment for variables (WA != NT)
- solutions are assignments satisfying all constraints

### Constraint Graph: node = variable; arc = constraint

### Backtracking Search

DFS + ordering variables + check constraints as you go

```
function RECURSIVE-BACKTRACKING(assignment, csp)
  if assignment is complete then return assignment
  var ← SELECT-UNASSIGNED-VARIABLE(Variables[csp], assignment, csp)
  for value in ORDER-DOMAIN-VALUES(var, assignment, csp) do
    if value is consistent with assignment given Constraints[csp] then
      add {var = value} to assignment
      result ← RECURSIVE-BACKTRACKING(assignment, csp)
      if result != failure then return result
      remove {var = value} from assignment
  return failure
```

### Ordering

- Minimum Remaining Values: choose the variable with the fewest legal left values in its domain
- Least Constraining Value: choose the value that rules out the fewest values in the remaining variables

### Filtering

- Forward Checking: cross off values in neighbor when assign value to node
- Constraint Propagation (AC-3): If X loses a value, neighbors of X need to be recheck consistency
- Arc Consistency: arc X > Y is consistent iff for every x in the tail there is some y could be assigned
- K-Consistency: For each k nodes, any consistent assignment to k-1 nodes we can find allowable value for the last node

### Tree-Structured CSP

- graph with no loop, can be solved in  $O(nd^2)$
- algorithm:
  - choose a root variable → order variables so that parents precede children
  - remove backward: for N:2, removeInconsistent(Parent(X), X)
  - assign forward: for 1:N, assign X with value consistent with Parent(X)
- cutset conditioning: init a set of variables such that the remaining constraint graph is a tree

### Min-Conflicts Algorithm

- a local search method: it's not optimal
- start from all variables assigned, while not solved, randomly select any conflicted variables, then reassign value that violates fewest constraints

## Local Search

- start from complete assignment (may not legal)
- reassign better value iteratively
- not optimal

### Hill-Climbing

start whatever, keep moving to best neighbor state until no neighbor is better

### Simulated Annealing

high temperature means more bad moves allowed, try to escape local minimum; temperature reduced by time

### Local Beam Search

run K local search simultaneously, select best K successors from all successors generated to be the new current states

## Bayes Net

- Joint Distribution: specify a real number for each assignment over a set of random variables
- Marginal Distribution: sub-tables of Joint with one random variable
- Conditional Probabilities:  $P(x|y) = P(x,y) / P(y)$
- Normalizing: scale the probabilities making the sum be 1

### Bayes' Rule

- $P(x|y) = P(y|x)P(x) / P(y)$
- two variables are independent if  $P(x, y) = P(x)P(y)$

### Bayes' Net

- a graph (DAG) to describe a set of random variables and their conditional independences
- Node: variables (with domains); can be observed or unobserved
- Arc: conditional independence
- CPT: describe node's conditional probability under each combination of its parent nodes

### Probabilistic Inference

- compute a probability of a query variable, given evidences
- Exact inference: sums of products of conditional probabilities from the BN

### Enumeration

- select entries consistent with the evidence
- sum out hidden variables
- normalize

### Variable Elimination

- pointwise product:
  - $P(B|A) \times P(A) = P(A,B)$
  - $P(U,V) \times P(V,W) \times P(W,X) = P(U,V,W,X)$
- summing out a variable
  - $\sum P(A,B) = P(A,b) + P(A, \neg b) = P(A)$
- interleave above steps

### Chains

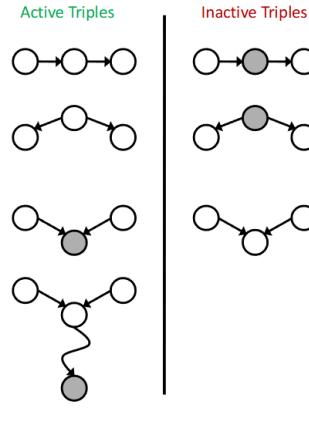
X indep from Z?	not given Y	given Y
Causal chains	X > Y > Z	No
Common causes	X < Y > Z	No
Common effect	X > Y < Z	Yes

### D-Separation

- X and Y conditionally independent given evidence variables {Z}?
- independence if no active paths (all inactive paths)
- active path if no inactive triples (all active triples)

### Markov Blanket

- As parents, childrens, and childrens' other parents
- If Markov blanket of A is observed, then A is conditional independent from all other variables



### Sampling

- a approximate inference method, used when 1) unknown distribution, 2) too expensive
- draw N samples from a sampling distribution S
- compute an approximate posterior probability, show this converges to the true probability P

### Prior Sampling

- e.g. we want Joint Distribution of (C, S, R, W)
  - sample xi from  $P(X_i | \text{Parent}(X_i)) \Rightarrow C$  from  $P(C)$ , S from  $P(S|C)$ , W from  $P(W|S,R)$
  - return  $(x_1, x_2, \dots, x_n) \Rightarrow (C=c, S=s, R=r, W=w)$
  - repeat the process to get a collection of  $(x_1, x_2, \dots, x_n)$

### Rejection Sampling

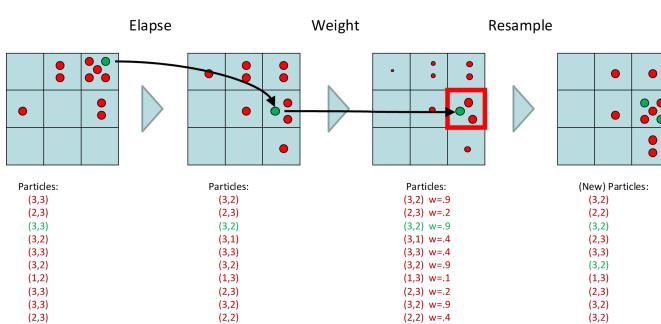
- e.g. we want  $P(C|+s)$ 
  - sample as in Prior Sampling
  - but ignore sample inconsistent with given evidences (i.e.  $S \neq +s$ )

### Likelihood Weighting

- for Rejection Sampling, we might reject too much sample if evidences are unlikely
- e.g. we want  $P(C|+s)$ 
  - sample as in Prior Sampling
  - but fix the value of evidence variables (keep  $S=+s$ )
  - and set the weight as  $w = w * P(x_i | \text{Parent}(X_i)) \Rightarrow w = w * P(+s|C)$
  - return  $(x_1, x_2, \dots, x_n)$ , w

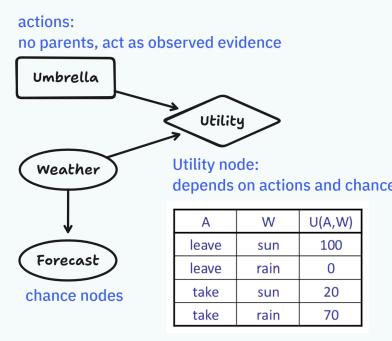
### Gibbs Sampling

- for Likelihood Weighting, evidence has no effect on upstream variables
- e.g. we want  $P(S|+r)$ 
  - fix evidence  $R=+r$
  - initialize other variables randomly
  - repeat:
    - choose one non-evidence variable X
    - resample X from  $P(X | \text{all other variables}) = P(X | \text{Markov}(X))$



## Decision Networks

- MEU: choose the action maximize the expected utility given the evidence
- bayes net with actions and utility



- Utility is defined by Weather and Umbrella
- assume the  $P(W=\text{sun}) = 0.7$
- we can calculate the expected utility of action
  - take:  $20 * 0.7 + 70 * 0.3 = 35$
  - leave:  $100 * 0.7 + 0 = 70 \Rightarrow$  optimal decision
- MEU( $\emptyset$ ) = 70
- Now consider if Forecast=bad
- $P(W|F) = P(F|W)P(W) / \sum_{w \in W} P(F|w)P(w)$
- $P(W|F=\text{bad}) = 0.34$
- $\text{MEU}(F=\text{bad}) = 53 \Rightarrow$  take

### Value of Information

- value of information = expect grow in MEU from this information
- example
  - without Forecast:  $\text{MEU}(\emptyset) = 70$
  - with Forecast:  $\text{MEU}(F=\text{bad}) = 53, \text{MEU}(F=\text{good}) = 95$
  - assume  $P(F=\text{bad}) = 0.59, P(F=\text{good}) = 0.41$
  - $\text{VPI}(F|\emptyset) = 0.59 * 53 + 0.41 * 95 - 70 = 7.8$

$$\text{VPI}(E'|e) = (\sum_{e'} P(e'|e)\text{MEU}(e, e')) - \text{MEU}(e)$$

- E: evidences we already know
- E': new information
- VPI is nonadditive, order-independent

## Markov Models

- value of X at a given time is called state
- transition model  $P(X_t | X_{t-1})$  specify how the state evolves over time
- transition probabilities are same all the time, future is independent of the past given present

### Markov Chain

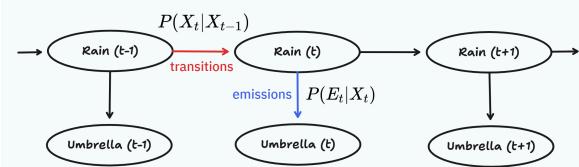
- $P(X_1 = \text{sun}) = 1$
- $P(X_2 = \text{sun}) = 0.9 * 1 + 0.3 * 0 = 0.9$
- Mini-Forward Algorithm: what's  $P(X)$  on some day t?

$$P(x_t) = \sum_{x_{t-1}} P(x_t | x_{t-1})P(x_{t-1})$$

- influence of initial state decrease by time
- the distribution ultimately converge in stationary distribution, where  $P_\infty(X) = P_{\infty+1}(X)$

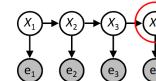
## Hidden Markov Models (HMM)

- hide Markov Chain over states X, we observe outputs at each time step
- future depends on past via the present
- current observation conditionally independent of all else given current state
- an HMM is defined by initial state + transitions + emissions

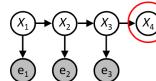


### Inference Tasks

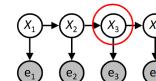
#### Filtering: $P(X_t | e_{1:t})$



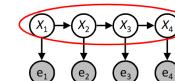
#### Prediction: $P(X_{t+k} | e_{1:t})$



#### Smoothing: $P(X_k | e_{1:t})$ , $k < t$



#### Explanation: $P(X_{1:t} | e_{1:t})$



### Filtering

- given all observations up to current, we want to "filter" out the distribution of current state  $B_t(X) = P(X_t | e_{1:t})$

### The Forward Algorithm

- passage of time: base on previous states to speculate current state
- observation: use new observation to adjust prediction on current state

$$P(x_t | e_{1:t}) \propto_{X_t} P(x_t, e_{1:t})$$

$$= \sum_{x_{t-1}} P(x_{t-1}, x_t, e_{1:t})$$

$$= \sum_{x_{t-1}} P(x_{t-1}, e_{1:t-1}) P(x_t | x_{t-1}) P(e_t | x_t)$$

$$= P(e_t | x_t) \sum_{x_{t-1}} P(x_{t-1}) P(x_t | x_{t-1}, e_{1:t-1})$$

### Particle Filtering

- representation: use N particle to represent possible current state
- elapse time: each particle move to new state by sample from transition model
- observe: if particle's new state more align to current observation, we give it higher weight
- resample: sample N new particle from weighted distribution